

Syllabus

CONCEPTUAL DATA MODELING

Database environment – Database system development lifecycle – Requirements collection – Database design - Entity-Relationship model – Enhanced-ER model – UML class diagrams.

RELATIONAL MODEL AND SQL

Relational model concepts - Integrity constraints - SQL Data manipulation – SQL Data definition – Views - SQL programming.

RELATIONAL DATABASE DESIGN AND NORMALIZATION

ER and EER-to-Relational mapping – Update anomalies – Functional dependencies – Inference rules – Minimal cover – Properties of relational decomposition – Normalization (upto BCNF).

TRANSACTIONS

Transaction Concepts – ACID Properties – Schedules – Serializability – Concurrency Control – Need for Concurrency – Locking Protocols – Two Phase Locking – Deadlock – Transaction Recovery - Save Points – Isolation Levels – SQL Facilities for Concurrency and Recovery.

NO-SQL DATABASES


No-SQL: CAP theorem – Document-based: MongoDB data model and CRUD operations; Column-based: Hbase data model and CRUD operations.

Unit V NO-SQL DATABASES


No-SQL: CAP theorem – Document-based: MongoDB data model and CRUD operations; Column-based: Hbase data model and CRUD operations.

SQL Databases

Relational




Analytical (OLAP)




NoSQL Databases

Key-Value



Column-Family



Graph




Document



- NoSQL databases use a variety of data models for accessing and managing data.
- These types of databases are optimized specifically for applications that require large data volume, low latency, and flexible data models
- Types of NoSQL Databases
 - Document Databases
 - Wide-Column Databases

Distributed Database

- A distributed database is a database that runs and stores data across multiple computers
- Example: NOSQL



NOSQL CAP THEOREM

Consistency


when system returns info, it is always up-to-date

Availability

system always returns info, even if stale

Partition Tolerance


system can continue operating during a partition



Replica Nodes: B, C

Fig: Distributed System

The CAP Theorem



Partition Happened;
Accept Or Pause Traffic?


Accept

Pause

Keep **Availability**,
Could Create
Consistency Problems

Keep **Consistency**,
But Lose **Availability**

- **A document database** is a type of NoSQL database that consists of sets of key-value pairs stored into a document



- Being a NoSQL database, you can easily store data without implementing a schema
- You can transfer the object model directly into a document using several different formats. The most commonly used are JSON, BSON, and XML.
- Examples of NoSQL document databases include MongoDB, CouchDB, Elasticsearch, and others.

Wide-Column Database

Wide-column stores are another type of NoSQL database. In them, data is stored and grouped into separately stored columns instead of rows. Such databases organize information into columns that function similarly to tables in relational databases.

Row-oriented			
ID	Name	Grade	GPA
001	John	Senior	4.00
002	Karen	Freshman	3.67
003	Bill	Junior	3.33

Column-oriented					
Name	ID	Grade	ID	GPA	ID
John	001	Senior	001	4.00	001
Karen	002	Freshman	002	3.67	002
Bill	003	Junior	003	3.33	003

- The most significant benefit of having column-oriented databases is that you can store large amounts of data within a single column
- Examples of popular wide-column databases include
 - HBase , Apache Cassandra, and CosmoDB.

MONGODB CRUD OPERATIONS

C → **Create**
R → **Read**
U → **Update**
D → **Delete**

DOCUMENT-ORIENTED DATABASE SCHEMA

Consider the **STUDENT** relation given below.

<u>SID</u>	SFName	SPhone
16s143	Ahmed	95214785
16j7890	Wafa	99663145

The above **STUDENT** relation is represented in **Document-Oriented database schema** as follows:

```
{_id: 16s143,  
SFName: "Ahmed",  
SPhone: 95214785  
}
```

Student_16s143 Document

```
{_id: 16j7890 ,  
SFName: "Wafa" ,  
SPhone: 99663145  
}
```

Convert the below **EMPLOYEE** relation into Document-Oriented Database Schema.

<u>EmpID</u>	EmpName	EmpAddress	EmpBdate
100	Laika Al-Mamari	Al Batinah, Sohar	1/21/1980
101	Khalid Al-Ameri	Al Aqur, Shinas	2/3/1990
102	Ranya Al-Balushi	Al Mutaqa, Sohar	5/25/1985

SOLUTION:

```
{_id: 100,  
  EmpName: "Laika Al-Mamari",  
  EmpAddress: "Al Batinah, Sohar",  
  EmpBdate: 1/21/1980  
}  
  
{_id: 101,  
  EmpName: "Khalid Al-Ameri",  
  EmpAddress: "Al Aqur, Shinas",  
  EmpBdate: 2/3/1990  
}  
  
{_id: 102,  
  EmpName: "Ranya Al-Balushi",  
  EmpAddress: "Al Mutaqa, Sohar",  
  EmpBdate: 5/25/1985  
}
```

Employee_100 document

Create Operations –

The create or insert operations are used to insert or add new documents in the collection. If a collection does not exist, then it will create a new collection in the database. You can perform, create operations using the following methods provided by the MongoDB:

Method	Description
<code>db.collection.insertOne()</code>	It is used to insert a single document in the collection.
<code>db.collection.insertMany()</code>	It is used to insert multiple documents in the collection.
<code>db.createCollection()</code>	It is used to create an empty collection.

```
[> use GeeksforGeeks
switched to db GeeksforGeeks
> db.student.insertOne({
... name : "Sumit",
... age : 20,
... branch : "CSE",
... course : "C++ STL",
... mode : "online",
... paid : true,
... amount : 1499
[... ]})
{
    "acknowledged" : true,
    "insertedId" : ObjectId("5e540cdc92e6dfa3fc48ddae")
}
> █
```

```
[> use GeeksforGeeks
switched to db GeeksforGeeks
> db.student.insertMany([
... {
...   name : "Sunit",
...   age : 20,
...   branch : "CSE",
...   course : "C++ STL",
...   mode : "online",
...   paid : true,
...   amount : 1499
... },
...
... {
...   name : "Rohit",
...   age : 21,
...   branch : "CSE",
...   course : "C++ STL",
...   mode : "online",
...   paid : true,
...   amount : 1499
... }
... ])
{
  "acknowledged" : true,
  "insertedIds" : [
    ObjectId("5e548d3192e6dfa3fc48ddaf"),
    ObjectId("5e548d3192e6dfa3fc48ddb0")
  ]
}
> █
```

- **Read Operations –**

- The Read operations are used to retrieve documents from the collection, or in other words, read operations are used to query a collection for a document.

<code>db.collection.find()</code>	It is used to retrieve documents from the collection.
-----------------------------------	---

`.pretty()`: this method is used to decorate the result such that it is easy to read.

Example : In this example, we are retrieving the details of students from the student collection using `db.collection.find()` method.

```
> use GeeksforGeeks
switched to db GeeksforGeeks
> db.student.find().pretty()
{
  "_id" : ObjectId("5e540cdc92e6dfa3fc48ddae"),
  "name" : "Sumit",
  "age" : 20,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499
}
{
  "_id" : ObjectId("5e540d3192e6dfa3fc48ddaf"),
  "name" : "Sumit",
  "age" : 20,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499
}
{
  "_id" : ObjectId("5e540d3192e6dfa3fc48ddb0"),
  "name" : "Rohit",
  "age" : 21,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499
}
>
```

Update Operations –

The update operations are used to update or modify the existing document in the collection. You can perform update operations using the following methods provided by the MongoDB:

Method	Description
<code>db.collection.updateOne()</code>	It is used to update a single document in the collection that satisfy the given criteria.
<code>db.collection.updateMany()</code>	It is used to update multiple documents in the collection that satisfy the given criteria.
<code>db.collection.replaceOne()</code>	It is used to replace single document in the collection that satisfy the given criteria.


```

> use GeeksforGeeks
switched to db GeeksforGeeks
> db.student.updateOne({name: "Sumit"},{$set:{age: 24 }})
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 0 }
> db.student.find().pretty()
{
  "_id" : ObjectId("5e540cdc92e6dfa3fc48ddae"),
  "name" : "Sumit",
  "age" : 24,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499
}
{
  "_id" : ObjectId("5e540d3192e6dfa3fc48ddaf"),
  "name" : "Sumit",
  "age" : 20,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499
}
{
  "_id" : ObjectId("5e540d3192e6dfa3fc48ddb0"),
  "name" : "Rohit",
  "age" : 21,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499
}
>

```

```
[> use GeeksforGeeks
switched to db GeeksforGeeks
[> db.student.updateMany({}, {$set: {year: 2020}})
{ "acknowledged" : true, "matchedCount" : 3, "modifiedCount" : 3 }
[> db.student.find().pretty()
{
  "_id" : ObjectId("5e540cdc92e6dfa3fc48ddae"),
  "name" : "Sumit",
  "age" : 24,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499,
  "year" : 2020
}
{
  "_id" : ObjectId("5e540d3192e6dfa3fc48ddaf"),
  "name" : "Sumit",
  "age" : 20,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499,
  "year" : 2020
}
{
  "_id" : ObjectId("5e540d3192e6dfa3fc48ddb0"),
  "name" : "Rohit",
  "age" : 21,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499,
  "year" : 2020
}
>
```

Delete Operations –

The delete operation are used to delete or remove the documents from a collection. You can perform delete operations using the following methods provided by the MongoDB:

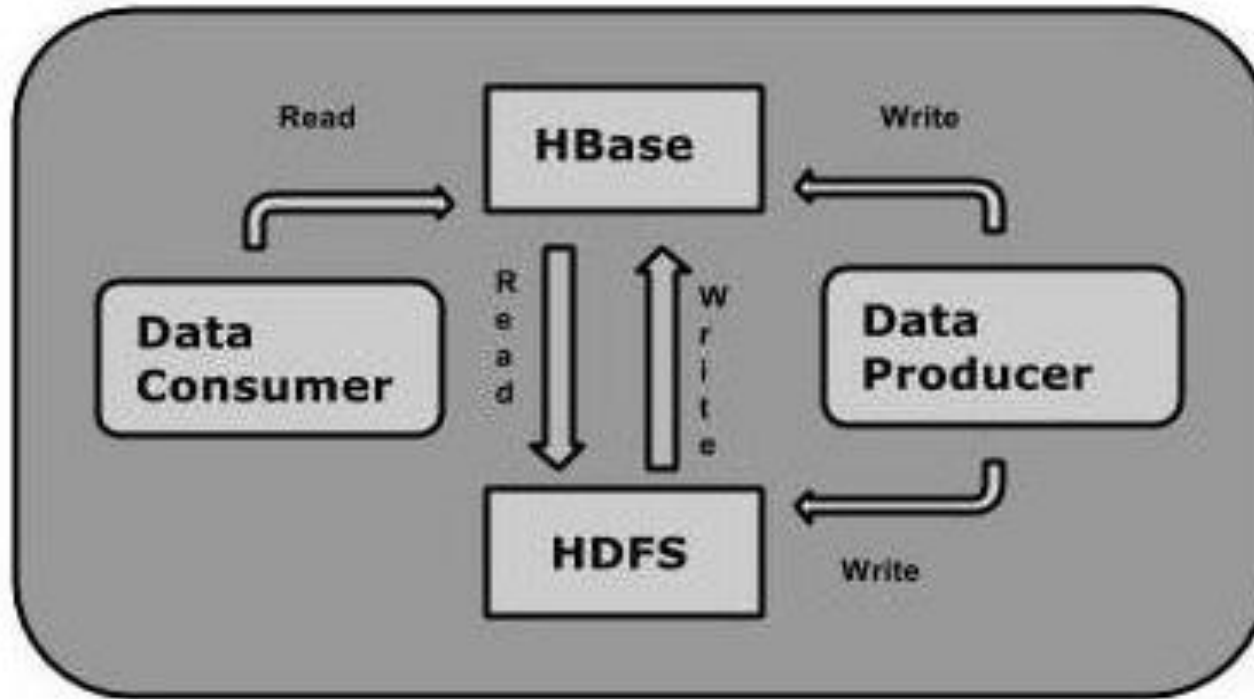
Method	Description
<code>db.collection.deleteOne()</code>	It is used to delete a single document from the collection that satisfy the given criteria.
<code>db.collection.deleteMany()</code>	It is used to delete multiple documents from the collection that satisfy the given criteria.

```
    "amount" : 1499
  }
> db.student.deleteOne({name: "Sunit"})
{ "acknowledged" : true, "deletedCount" : 1 }
> db.student.find().pretty()
{
  "_id" : ObjectId("5e548d3192e6dfa3fc48ddaf"),
  "name" : "Sunit",
  "age" : 20,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499,
  "year" : 2020
}
{
  "_id" : ObjectId("5e54103592e6dfa3fc48ddb1"),
  "name" : "Rohit",
  "age" : 21,
  "branch" : "CSE",
  "course" : "C++ STL",
  "mode" : "online",
  "paid" : true,
  "amount" : 1499
}
```

```
> db.student.deleteMany({})  
{ "acknowledged" : true, "deletedCount" : 2 }  
> █
```

HBASE

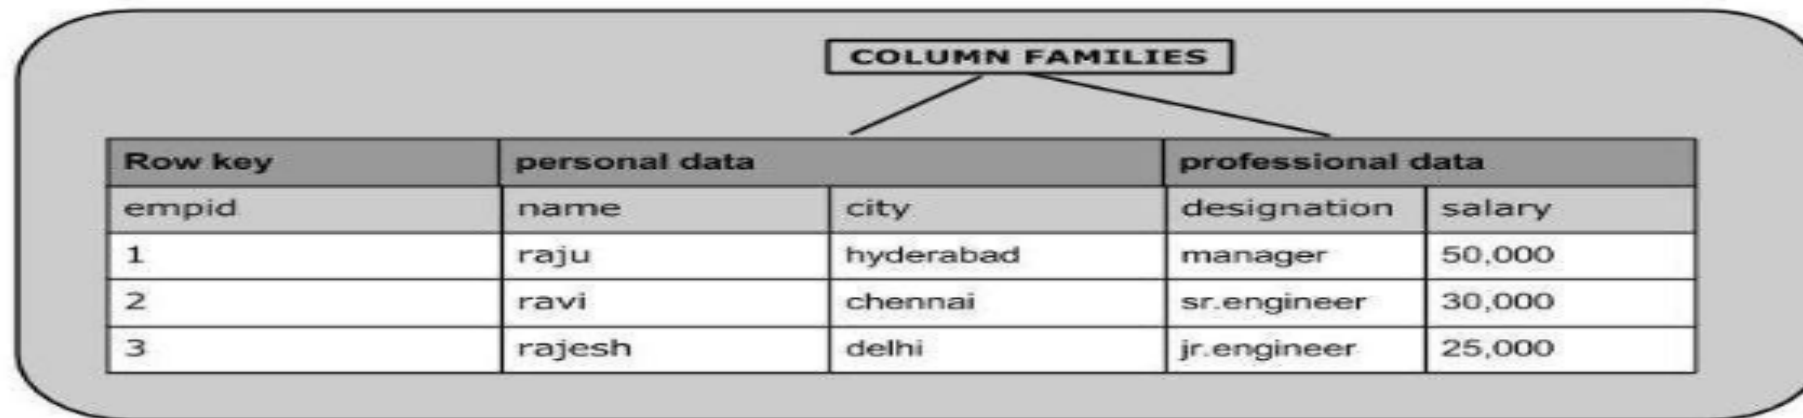
- HBase is a distributed column-oriented database built on top of the Hadoop file system. It is an open-source project and is horizontally scalable.
- it leverages the fault tolerance provided by the Hadoop File System (HDFS).



HDFS	HBase
HDFS is a distributed file system suitable for storing large files.	HBase is a database built on top of the HDFS.
HDFS does not support fast individual record lookups.	HBase provides fast lookups for larger tables.
It provides high latency batch processing; no concept of batch processing.	It provides low latency access to single rows from billions of records (Random access).
It provides only sequential access of data.	HBase internally uses Hash tables and provides random access, and it stores the data in indexed HDFS files for faster lookups.

Row-Oriented Database	Column-Oriented Database
It is suitable for Online Transaction Process (OLTP).	It is suitable for Online Analytical Processing (OLAP).
Such databases are designed for small number of rows and columns.	Column-oriented databases are designed for huge tables.

The following image shows column families in a column-oriented database:



HBase	RDBMS
HBase is schema-less, it doesn't have the concept of fixed columns schema; defines only column families.	An RDBMS is governed by its schema, which describes the whole structure of tables.
It is built for wide tables. HBase is horizontally scalable.	It is thin and built for small tables. Hard to scale.
No transactions are there in HBase.	RDBMS is transactional.
It has de-normalized data.	It will have normalized data.
It is good for semi-structured as well as structured data.	It is good for structured data.

- Structure of HBASE table

	Row Key	Family "Details"	Family "Relatives"	Family "Accounts"
Row	<value>	Name Nickname Address Email Mobile Fax Home	Wife Sister	Checking Savings Business
Row	<value>	Name Address Email Mobile		Checking
Row	<value>	Name Nickname Address Email Home	Father Mother	Savings

	Row Key	Family "Details"	Family "Relatives"
Row	101	Name: Adam Mobile: 555-555-1234 Nickname: A-Man Fax: 555-555-2222 Address: 123 Main Home: 555-234-5325 Email: Adam@email.com	Wife: Debby Sister: Kim
Row	102	Name: Bob Address: 12 East St. Email: Bob@email.com Mobile: 555-562-1234	
Row	103	Name: Christopher Home: 555-232-3332 Nickname: Chris Address: 504 Rogers Road Email: Chris@email.com	Father: Thomas Mother: Casey

Creating a table in HBase - create

	Row Key	Family "Details"	Family "Relatives"	Family "Accounts"
Row	101	Name: Adam Nickname: A-Man Address: 123 Main Email: Adam@email.com Mobile: 555-555-1234 Fax: 555-555-2222 Home: 555-234-5325	Wife: Debby Sister: Kim	Checking: \$1,500 Savings: \$25,000 Business: \$8,250
Row	102	Name: Bob Address: 12 East St. Email: Bob@email.com Mobile: 555-562-1234		Checking: \$250
Row	103	Name: Christopher Nickname: Chris Address: 504 Rogers Road Email: Chris@email.com Home: 555-232-3332	Father: Thomas Mother: Casey	Savings: \$2,000

- General format of: `create <table>, <CF1>, <CF2>, ... <CFn>`
`create 'customers', 'details', 'relatives', 'accounts'`

Inserting and updating data in HBase - put

	Row Key	Family "Details"	Family "Relatives"	Family "Accounts"
Row	101	Name: Adam Nickname: A-Man Address: 123 Main Email: Adam@email.com Mobile: 555-555-1234 Fax: 555-555-2222 Home: 555-234-5325	Wife: Debby Sister: Kim	Checking: \$1,500 Savings: \$25,000 Business: \$8,250
Row	102	Name: Bob Address: 12 East St. Email: Bob@email.com Mobile: 555-562-1234		Checking: \$250
Row	103	Name: Christopher Nickname: Chris Address: 504 Rogers Road Email: Chris@email.com Home: 555-232-3332	Father: Thomas Mother: Casey	Savings: \$2,000

- General format of: put <table>, <row key>, <CF:Qualifier>, <value>
put 'customers', '101', 'details:name', 'Adam'

Deleting data

- To delete the value of a single family qualifier:
 - Delete <table>, <row>, <column family>, <qualifier>, <timestamp>
delete 'customers', '101', 'accounts:business'
- To delete an entire row:
 - Deleteall <table> , <row>
deleteall 'customers', '102'
- To drop a table:
 - First disable the table, then drop it
disable 'customers'
drop 'customers'

- Update table

```
hbase(main):002:0> scan 'customer'
```

```
ROW
```

```
1  
1  
1  
1
```

```
1 row(s) in 0.1880 seconds
```

```
COLUMN+CELL
```

```
column=customer_contact_details:email, timestamp=1667405992288, value=rahulroy@gmail.com  
column=customer_contact_details:mobile, timestamp=1667405934134, value=9999767767  
column=customer_info:first_name, timestamp=1667405716754, value=Rahul  
column=customer_info:last_name, timestamp=1667405770074, value=Roy
```

```
hbase(main):003:0> get 'customer','1'
```

```
COLUMN
```

```
customer_contact_details:email  
customer_contact_details:mobile  
customer_info:first_name  
customer_info:last_name
```

```
4 row(s) in 0.0200 seconds
```

```
CELL
```

```
timestamp=1667405992288, value=rahulroy@gmail.com  
timestamp=1667405934134, value=9999767767  
timestamp=1667405716754, value=Rahul  
timestamp=1667405770074, value=Roy
```

```
hbase(main):004:0> put 'customer','1','customer_contact_details:mobile','8888999000'
```

```
0 row(s) in 0.0830 seconds
```

```
hbase(main):005:0> scan 'customer'
```

```
ROW
```

```
1  
1  
1  
1
```

```
1 row(s) in 0.0310 seconds
```

```
COLUMN+CELL
```

```
column=customer_contact_details:email, timestamp=1667405992288, value=rahulroy@gmail.com  
column=customer_contact_details:mobile, timestamp=1667407951467, value=8888999000  
column=customer_info:first_name, timestamp=1667405716754, value=Rahul  
column=customer_info:last_name, timestamp=1667405770074, value=Roy
```

```
hbase(main):006:0> █
```



```
hbase(main):004:0> create 'Customers', 'details', 'relatives', 'accounts'
0 row(s) in 1.2390 seconds

=> Hbase::Table - Customers
hbase(main):005:0> list
TABLE
Customers
customers
2 row(s) in 0.0090 seconds

=> ["Customers", "customers"]
hbase(main):006:0> disable 'Customers'
0 row(s) in 2.2470 seconds

hbase(main):007:0> drop 'Customers'
0 row(s) in 1.2330 seconds

hbase(main):008:0> put 'customers', '101', 'details:name', 'Adam'
0 row(s) in 0.1440 seconds

hbase(main):009:0> scan 'customers'
ROW                                COLUMN+CELL
 101                                column=details:name, timestamp=1599233344525, va
1 row(s) in 0.0300 seconds

hbase(main):010:0> █
```

- Get- used to display specific row of table
- Syntax: get 'table_name','row_id'

```
hbase(main):043:0> get 'customers', '102'
COLUMN                                CELL
accounts:checking                    timestamp=1599233413328, value=250
details:address                      timestamp=1599233413249, value=12 East St.
details:email                        timestamp=1599233413269, value=bob@email.com
details:mobile                       timestamp=1599233413300, value=555-562-1234
details:name                         timestamp=1599233413227, value=Bob
1 row(s) in 0.0210 seconds
```

```

hbase(main):043:0> get 'customers', '102'
COLUMN                                CELL
accounts:checking                     timestamp=1599233413328, value=250
details:address                       timestamp=1599233413249, value=12 East St.
details:email                         timestamp=1599233413269, value=bob@email.com
details:mobile                        timestamp=1599233413300, value=555-562-1234
details:name                          timestamp=1599233413227, value=Bob
1 row(s) in 0.0210 seconds

hbase(main):044:0> get 'customers', '102', 'details'
COLUMN                                CELL
details:address                       timestamp=1599233413249, value=12 East St.
details:email                         timestamp=1599233413269, value=bob@email.com
details:mobile                        timestamp=1599233413300, value=555-562-1234
details:name                          timestamp=1599233413227, value=Bob
1 row(s) in 0.0170 seconds

hbase(main):045:0> get 'customers', '102', 'details', 'accounts'
COLUMN                                CELL
accounts:checking                     timestamp=1599233413328, value=250
details:address                       timestamp=1599233413249, value=12 East St.
details:email                         timestamp=1599233413269, value=bob@email.com
details:mobile                        timestamp=1599233413300, value=555-562-1234
details:name                          timestamp=1599233413227, value=Bob
1 row(s) in 0.0130 seconds

hbase(main):046:0> get 'customers', '102', {COLUMNS => 'details:name', 'accounts'

```

- List- display all list of table available

```
hbase(main):027:0> list
TABLE
ns_db1.table1
t
2 row(s) in 0.0480 seconds

=> ["ns_db1.table1", "t"]
hbase(main):028:0> █
```


- Create – used to create a new table
- Syntax: create 'table_name', 'column_name',...'column_name'

```
hbase(main):028:0> create 'employee', 'personal data', 'professional data'
0 row(s) in 0.8790 seconds

=> Hbase::Table - employee
hbase(main):029:0> list
TABLE
employee
ns_db1.table1
t
3 row(s) in 0.0150 seconds

=> ["employee", "ns_db1.table1", "t"]
hbase(main):030:0> 
```

- Put- Used to insert a record into the tables
- Syntax: put 'table_name', 'id','column_name:attribute','value'
- Scan- used to display table data
- Syntax: scan 'table name'

```
hbase(main):030:0> put 'employee', '1', 'personal data:name', 'ajay'
0 row(s) in 0.0710 seconds

hbase(main):031:0> scan 'employee'
ROW                                COLUMN+CELL
 1                                column=personal data:name, timestamp=1479408494256, value=ajay
1 row(s) in 0.0890 seconds

hbase(main):032:0> put 'employee', '1', 'professional data:designation', 'Manager'
0 row(s) in 0.0720 seconds

hbase(main):033:0> scan 'employee'
ROW                                COLUMN+CELL
 1                                column=personal data:name, timestamp=1479408494256, value=ajay
 1                                column=professional data:designation, timestamp=1479408637299, value=Manager
1 row(s) in 0.1040 seconds
```

1 row(s) in 0.1040 seconds

```
hbase(main):034:0> put 'employee', '2', 'personal data:name', 'sam'
0 row(s) in 0.0280 seconds
```

```
hbase(main):035:0> scan 'employee'
ROW                                COLUMN+CELL
 1                                column=personal data:name, timestamp=1479408494256, value=ajay
 1                                column=professional data:designation, timestamp=1479408637299, value=Manager
 2                                column=personal data:name, timestamp=1479408704732, value=sam
2 row(s) in 0.0900 seconds
```

```
hbase(main):036:0> put 'employee', '2', 'professional data:designation', 'QA'
0 row(s) in 0.0260 seconds
```

```
hbase(main):037:0> █
```

\$ hbase shell

- Getting the status of the system and number of servers

hbase(main):002:0> status

- Creating a table:

hbase(main):005:0> create 'emp', 'personal details', 'professional details'

- Describe the table:

hbase(main):017:0> describe 'emp'

- List the tables present in keyspace:

hbase(main):001:0> list

➤ Inserting data into the table:

```
hbase(main):018:0> put 'emp','1','personal details:name','Ram'
```

➤ Viewing records inserted in the table:

```
hbase(main):023:0> scan 'emp'
```

➤ Getting the record from Hbase table:

```
hbase(main):026:0> get 'emp', '1'
```

➤ Getting a specific column from the record:

```
hbase(main):002:0> get 'emp', '1', { COLUMN => 'personal details:name'}
```

➤ Dropping a table. We can drop a table by first disabling it and then executing the dropped table:

```
hbase(main):016:0> disable 'emp'
```

```
hbase(main):017:0> drop 'emp'
```